

BCNF

Definition: No data redundancy in BCNF

Consider a relation schema R with FDs F

R is in Boyce-Codd Normal Form (BCNF) if for every non-trivial FD $a \rightarrow A$ therefore A is definitely non-prime attributes

a is a superkey of R & non-prime attributes are dependent on superkey

Violation: A non-trivial FD $a \rightarrow b$ that holds on R is said to violate BCNF if

a is not a superkey of R

Decomposition:

A decomposition $\{R_1, R_2, \dots, R_n\}$ of R (with FDs F) is in BCNF if each R_i is in BCNF (w.r.t. $F|_{R_i}$). No need to refer to original F

Algorithm 5: BCNF Test for R_i look for FD that violate BCNF

Input: F is a set of FDs that hold on a schema R

Output: A completely non-trivial FD that violates BCNF if R_i is not in BCNF; otherwise, $null$

1. if (R_i has exactly 2 attributes) then return $null$

2. for each $(a \subseteq R_i$ such that $a \neq \emptyset$)

3. let $X = a^*$ in R_i (w.r.t. F)

// $a \rightarrow X$ is an FD that holds on R_i

4. if ($a \subset X \subset R_i$) then return $a \rightarrow X - a$

5. return $null$

Algorithm 6: BCNF Decomposition

Input: Schema R with FDs F

Output: A BCNF decomposition of R

1. initialize $\delta = \emptyset$; $i = 1$; $\text{temp} = \{R\}$

2. while ($\text{temp} \neq \emptyset$)

3. remove some R' from temp

4. $f = \text{BCNFTest}(F, R')$ // Algorithm 5 look for FD that violate BCNF

5. if ($f = null$) then

6. $\delta = \delta \cup \{R'\}$

7. else

There is a violation of BCNF. Here make sure it doesn't violate BCNF.

let f be $a \rightarrow b$

let $c = R' - b$

temp = temp $\cup \{R_i(ab), R_{i+1}(c)\}$ // for next check

11. $i = i + 2$

guaranteed to have lossless join decomposition becos its decomposed based on FD

12. return δ

Properties:

Not guaranteed losslessjoin unless algo6 used

BCNF decompositions are lossless-join decompositions

However, not all schema has a dependency-preserving BCNF decompositions

BCNF decompositions are not necessarily dependency-preserving

3NF

Consider a relation schema R with FDs F

R is in Third Normal Form (3NF) if for every non-trivial FD $a \rightarrow A$ in F

a is a superkey of R OR // BCNF

A is a prime attribute // but less restrictive

Violation: BCNF guarantees no data redundancy, but not all schema has BCNF

Therefore 3NF is more relax. If a is not superkey, we allow A to be prime attribute

A non-trivial FD $a \rightarrow b$ that holds on R is said to violate BCNF if

a is not a superkey of R AND

A is a non-prime attribute

Decomposition:

A decomposition $\{R_1, R_2, \dots, R_n\}$ of R (with FDs F) is in 3NF if each R_i is in 3NF (w.r.t. $F|_{R_i}$)

3NF property: All schema can be decompose to 3NF, 3NF is not cfm lossless join, 3NF not cfm dependency preserving, some data may be redundant (as compared to BCNF definitely no redundant data)

Algorithm 7: 3NF Decomposition

Input: Schema R with FDs F which is a minimal cover

Output: A dependency-preserving and lossless-join 3NF decomposition of R

1. initialize $\delta = \emptyset$

2. apply rule 0 to combine FDs in F with same L.H.S into a single FD

3. let $G = \{f_1, f_2, \dots, f_n\}$ be the resultant set of FDs

4. for each (FD f_i of the form $a_i \rightarrow b_i$ in G)

5. create a relation schema $R_i(a_i, b_i)$ for FD f_i

6. insert the relation schema $R_i(a_i, b_i)$ into δ

7. choose a key K of R and insert a relation schema $R_{n+1}(K)$ into δ

8. remove redundant relation schema from δ as follows:

9. delete R_i from δ if $\exists R_j \in \delta : i \neq j \wedge R_i \subseteq R_j$

10. return δ

Algorithm 7: Properties

Each $R_i \in \delta$ is in 3NF

δ is a dependency-preserving decomposition

δ is a lossless-join decomposition

The decomposition δ produced by Algorithm 7 (i.e., synthesis approach) may not be unique

Depends on choice of minimal cover

Depends on choice of redundant relation schema being removed

SNF is not unique

ER

0 or more: Partial participation constraint

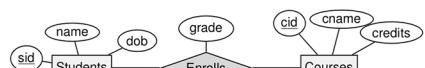
Each instance of E participates in at most one instance of R Key constraint

Each instance of E participates in at least one instance of R Total participation constraint

Each instance of E participates in exactly one instance of R

E is a weak entity set with identifying owner E' & identifying relationship set R

Relationship Sets w/o Constraints



Many-to-many: just take all primary key and stick them tgt

```

create table Enrolls (
  sid char(20) references Students,
  cid char(5),
  grade numeric,
  primary key (sid, cid),
  foreign key (cid) references Courses;
  
```

Relationship Sets with Key Constraints



Better approach

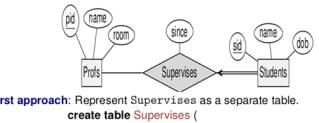
First approach: Represent Supervises using a separate table

- Pros (pid, name, room)
- Students (sid, name, dob)
- Supervises (sid, pid, since)

Second approach: Combine Supervises & Students into one table Works becos we can just assign null to student if student dont present

- Pros (pid, name, room)
- SupervisedStudents (sid, name, dob, pid, since)

Key & Total Participation Constraints



First approach: Represent Supervises as a separate table.

```

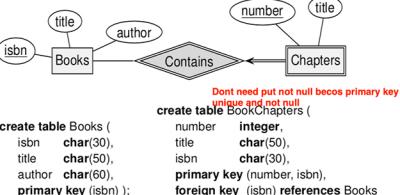
create table Supervises (
  sid char(20) primary key,
  pid char(7),
  since date,
  foreign key (sid) references Students,
  foreign key (pid) references Pros;
  
```

Total participation constraint of Students w.r.t. Supervises is not captured by database schema!

Means I can insert a new record to Student, but don't insert a new record to Supervises. The previous one ok becos it is at most one, means can't insert. But this is exactly one, must insert

Weak Entity Sets

Weak entity set & its identifying relationship set are represented by a single relation



Dont need put not null becos primary key is unique and not null

```

create table Books (
  isbn char(30),
  title char(50),
  author char(60),
  primary key (isbn));
create table BookChapters (
  number integer,
  title char(50),
  isbn char(30),
  primary key (number, isbn),
  foreign key (isbn) references Books
  on delete cascade);
  
```

Table is subset of another

where not exists (

```

select 1 from
  (select Works.eid from Works
  where Works.pid = Projects.pid) as T1
  left outer join
  (select Specializes.eid from Specializes
  where Specializes.aid = 'A') as T2
  on T1.eid = T2.eid
  where T2.eid is null
  )
```

exist / not exists

Returns true if the result subquery is non-empty; otherwise, false

where exists (

```

select 1
from Sells S
where S.rname = 'Corleone Corner'
and S.pizza = L.pizza
);
```

if you dont use 'exists':

case

select name, case

```

when marks >= 70 then 'A'
when marks >= 60 then 'B'
when marks >= 50 then 'C'
else 'D'
end as grade
```

from Scores;

select name, case

```

when (first = 'pass') or (second = 'pass')
or (third = 'pass') then 'pass'
else 'fail'
end as result
from Tests;
```

like

if '_____e' means exactly 4 chars ending with e
Can stick the % anywhere, it means 0 or more characters

coalesce

select name, coalesce(third,second,first) as result
from Tests;

coalesce at least 2 argument. If all null, value is null

Can also use: coalesce(argument1, 'explicit value')

- coalesce returns the first non-null value in its arguments
- Null is returned if all arguments are null

Relationship Sets with Key Constraints (cont.)

Cannot make relationship attribute as key
The attribute must be present in some entity



First approach: create table Supervises (

```

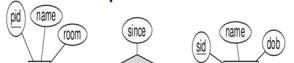
sid char(20),
pid char(7),
since date,
primary key (sid),
foreign key (sid) references Students,
foreign key (pid) references Pros);
  
```

Second approach: create table SupervisedStudents (

```

sid char(20),
name char(30),
dob date,
pid char(7),
since date,
primary key (sid),
foreign key (pid) references Pros);
  
```

Key & Total Participation Constraints



Better approach

Second approach: Combine Students & SupervisedStudents into a single table create table SupervisedStudents (

```

sid char(20),
name char(30),
dob date,
pid char(7),
since date,
primary key (pid),
foreign key (pid) references Pros);
  
```

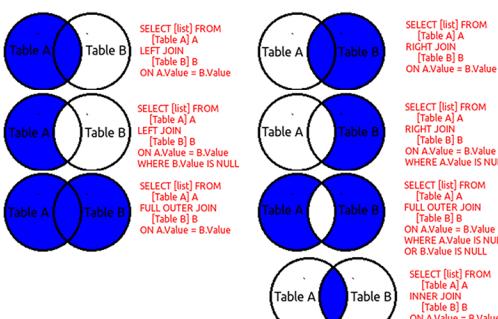
IS-A

- Overlap constraint: Can an entity belong to multiple subclasses?

• A ISA hierarchy satisfies the overlap constraint if an entity in a superclass could belong to multiple subclasses

- Covering constraint Does an entity in a superclass have to belong to some subclass?

• A ISA hierarchy satisfies the covering constraint if every entity in a superclass has to belong to some subclass



in

Subquery must return exactly one column
Returns false if result of subquery is empty; otherwise return the result of the boolean expression

((v = v₁) or (v = v₂) or ... or (v = v_n))

where

v denote the result of expression

{v₁, v₂, ..., v_n} denote the result of subquery

where

pizza in (

select pizza this subquery must return only 1 column

otherwise its not meaningful to compare

);

any

Subquery must return exactly one column

Returns false if result of subquery is empty; otherwise return the result of the boolean expression

The operator must be a standard comparison operator (=, <, >, <=, >=, or <>).

((v op v₁) or (v op v₂) or ... or (v op v_n))

select distinct name

from Sells

where rname < 'Corleone Corner'

price > any (/all

select price

from Sells

where rname = 'Corleone Corner'

);

nullif

select name, nullif(result,'absent') as stat

from Tests;

• nullif (value₁, value₂)

• Returns null if value₁ is equal to value₂; otherwise returns value₁

table_name ADD COLUMN new_column_name TYPE;

table_name RENAME COLUMN column_name;

ALTER TABLE table_name ALTER COLUMN column_name SET DEFAULT value | DROP DEFAULT;

ALTER TABLE table_name ALTER COLUMN column_name [SET NOT NULL | DROP NOT NULL];

ALTER TABLE table_name ADD CHECK expression;

ALTER TABLE table_name RENAME TO new_table_name;

ALTER TABLE table_name R